

## What is Classification?

- Supervised Learning is divided into regression and classification where regression is used when the target variable is continuous and classification is used when the target variable is discrete.

### Classification Techniques

- Logistic Regression
- Support Vector Machines
- Principal Component Analysis
- Decision Tree
- Ensemble Techniques
- K Nearest Neighbors
- Naive Bayes

#### 1. Logistic Regression

- Classification Algorithm that classifies based on probability.
- Uses Sigmoid curve as a cost function.
- We use sigmoid function to map predicted values to probabilities between 1 and 0.
- A sigmoid function is represented as below

$$f(x) = \frac{1}{1+e^{-x}}$$


```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import math
import os

from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, auc, confusion_matrix, roc_auc_score, roc_curve, recall_score
```

```
In [2]: array = np.arange(-6,6,0.1)
sigmoid = []
for i in array:
    sig = 1/(1+math.exp(-i))
    sigmoid.append(sig)
```

```
In [3]: plt.plot(array,sigmoid)
plt.show()
```



```
In [4]: df = pd.read_csv('./bank.csv', sep = ';')
df.head()
```

Out[4]:

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campaign	pdays	previous	poutcome	
0	30	unemployed	married	primary	no	1787	no	no	cellular	19	oct	79	1	-1	0	unknown	n
1	33	services	married	secondary	no	4789	yes	yes	cellular	11	may	220	1	339	4	failure	n
2	35	management	single	tertiary	no	1350	yes	no	cellular	16	apr	185	1	330	1	failure	n
3	30	management	married	tertiary	no	1476	yes	yes	unknown	3	jun	199	4	-1	0	unknown	n
4	59	blue-collar	married	secondary	no	0	yes	no	unknown	5	may	226	1	-1	0	unknown	n

The classification goal is to predict if the client will subscribe (Yes/No) a term deposit (variable y).

## Data

```
In [5]: df.shape
```

Out[5]: (4521, 17)

```
In [6]: df.dtypes[df.dtypes == 'object'].index
```

Out[6]: Index(['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', 'month', 'poutcome', 'y'], dtype='object')

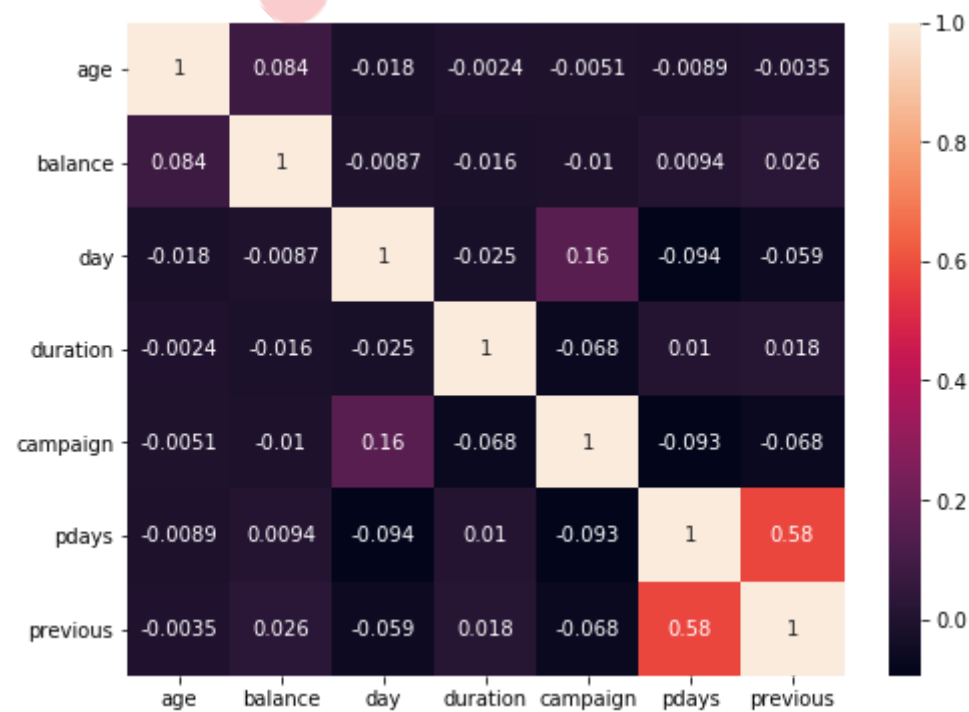
## 4. Checking Multi-Collinearity

```
In [7]: corr = df.corr()
corr
```

Out[7]:

	age	balance	day	duration	campaign	pdays	previous
age	1.000000	0.083820	-0.017853	-0.002367	-0.005148	-0.008894	-0.003511
balance	0.083820	1.000000	-0.008677	-0.015950	-0.009976	0.009437	0.026196
day	-0.017853	-0.008677	1.000000	-0.024629	0.160706	-0.094352	-0.059114
duration	-0.002367	-0.015950	-0.024629	1.000000	-0.068382	0.010380	0.018080
campaign	-0.005148	-0.009976	0.160706	-0.068382	1.000000	-0.093137	-0.067833
pdays	-0.008894	0.009437	-0.094352	0.010380	-0.093137	1.000000	0.577562
previous	-0.003511	0.026196	-0.059114	0.018080	-0.067833	0.577562	1.000000

```
In [8]: plt.figure(figsize=(8,6))
sns.heatmap(corr,annot=True)
plt.show()
```



```
In [9]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4521 entries, 0 to 4520
Data columns (total 17 columns):
age                4521 non-null int64
job                4521 non-null object
marital            4521 non-null object
education          4521 non-null object
default            4521 non-null object
balance            4521 non-null int64
housing            4521 non-null object
loan               4521 non-null object
contact            4521 non-null object
day                4521 non-null int64
month              4521 non-null object
duration           4521 non-null int64
campaign           4521 non-null int64
pdays             4521 non-null int64
previous           4521 non-null int64
poutcome           4521 non-null object
y                  4521 non-null object
dtypes: int64(7), object(10)
memory usage: 600.6+ KB
```

```
In [10]: df.dtypes[df.dtypes == 'int64'].index
```

```
Out[10]: Index(['age', 'balance', 'day', 'duration', 'campaign', 'pdays', 'previous'], dtype='object')
```

```
In [11]: # consider numerical data for calculating Variance Inflation Factor
num_col = df.dtypes[df.dtypes == 'int64'].index
```

```
In [12]: X_num = df[num_col]
X_num.shape
```

```
Out[12]: (4521, 7)
```

```
In [13]: from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
In [14]: vif = pd.DataFrame()
vif['Features'] = X_num.keys()
vif['Values'] = [variance_inflation_factor(X_num.values,i) for i in range(7)]
vif
```

```
Out[14]:
```

	Features	Values
0	age	5.092604
1	balance	1.231819
2	day	4.057535
3	duration	1.928720
4	campaign	1.830360
5	pdays	1.733904
6	previous	1.655651

There is no multi-collinearity effect in the dataset since all VIF values are less than ~5

## Create Dummy Features

```
In [15]: df_dum = pd.get_dummies(df,drop_first = True)
```

```
In [16]: df_dum.head()
```

```
Out[16]:
```

	age	balance	day	duration	campaign	pdays	previous	job_blue-collar	job_entrepreneur	job_housemaid	...	month_jun	month_mar	month_may	n
0	30	1787	19	79	1	-1	0	0	0	0	...	0	0	0	
1	33	4789	11	220	1	339	4	0	0	0	...	0	0	1	
2	35	1350	16	185	1	330	1	0	0	0	...	0	0	0	
3	30	1476	3	199	4	-1	0	0	0	0	...	1	0	0	
4	59	0	5	226	1	-1	0	1	0	0	...	0	0	1	

5 rows × 43 columns

## Feature Engineering

```
In [17]: # Splitting data into independent and dependent
X = df_dum.iloc[:, :-1] # independent variables
y = df_dum.iloc[:, -1] # dependent variables
```

```
In [18]: y.value_counts()/len(y)
```

```
Out[18]: 0    0.88476
         1    0.11524
         Name: y_yes, dtype: float64
```

There is clearly unbalance in the dataset

```
In [19]: import statsmodels.api as sm
```



```
In [20]: model = sm.GLM(y,X).fit()
print(model.summary())
```

```

Generalized Linear Model Regression Results
=====
Dep. Variable:          y_yes      No. Observations:          4521
Model:                  GLM        Df Residuals:              4479
Model Family:           Gaussian   Df Model:                  41
Link Function:          identity    Scale:                    0.073097
Method:                  IRLS      Log-Likelihood:           -480.52
Date:                   Sat, 22 Feb 2020    Deviance:                 327.40
Time:                   12:52:00    Pearson chi2:             327.
No. Iterations:         3
Covariance Type:        nonrobust
=====


```

	coef	std err	z	P> z	[0.025	0.975]
age	0.0005	0.000	1.267	0.205	-0.000	0.001
balance	-7.02e-07	1.38e-06	-0.510	0.610	-3.4e-06	2e-06
day	0.0017	0.001	3.156	0.002	0.001	0.003
duration	0.0005	1.56e-05	31.097	0.000	0.000	0.001
campaign	-0.0015	0.001	-1.069	0.285	-0.004	0.001
pdays	7.641e-06	7.52e-05	0.102	0.919	-0.000	0.000
previous	0.0005	0.003	0.152	0.879	-0.006	0.007
job_blue-collar	-0.0224	0.015	-1.453	0.146	-0.053	0.008
job_entrepreneur	-0.0151	0.025	-0.608	0.543	-0.064	0.034
job_housemaid	-0.0249	0.029	-0.854	0.393	-0.082	0.032
job_management	-0.0042	0.017	-0.241	0.809	-0.038	0.030
job_retired	0.0554	0.024	2.290	0.022	0.008	0.103
job_self-employed	-0.0065	0.024	-0.271	0.786	-0.054	0.041
job_services	-0.0075	0.018	-0.418	0.676	-0.043	0.028
job_student	0.0578	0.033	1.761	0.078	-0.007	0.122
job_technician	-0.0146	0.016	-0.925	0.355	-0.046	0.016
job_unemployed	-0.0430	0.027	-1.583	0.113	-0.096	0.010
job_unknown	0.0454	0.047	0.971	0.332	-0.046	0.137
marital_married	-0.0342	0.012	-2.738	0.006	-0.059	-0.010
marital_single	-0.0140	0.014	-0.993	0.321	-0.042	0.014
education_secondary	0.0074	0.012	0.598	0.550	-0.017	0.032
education_tertiary	0.0270	0.015	1.756	0.079	-0.003	0.057
education_unknown	-0.0253	0.023	-1.102	0.271	-0.070	0.020
default_yes	0.0495	0.032	1.568	0.117	-0.012	0.111
housing_yes	-0.0153	0.010	-1.613	0.107	-0.034	0.003
loan_yes	-0.0330	0.012	-2.867	0.004	-0.056	-0.010
contact_telephone	0.0039	0.017	0.231	0.817	-0.029	0.037
contact_unknown	-0.0776	0.014	-5.613	0.000	-0.105	-0.051
month_aug	-0.0256	0.020	-1.285	0.199	-0.065	0.013
month_dec	0.0578	0.063	0.917	0.359	-0.066	0.181
month_feb	0.0275	0.024	1.141	0.254	-0.020	0.075
month_jan	-0.0959	0.028	-3.459	0.001	-0.150	-0.042
month_jul	-0.0571	0.019	-3.000	0.003	-0.094	-0.020
month_jun	0.0404	0.023	1.792	0.073	-0.004	0.084
month_mar	0.2289	0.042	5.456	0.000	0.147	0.311
month_may	-0.0267	0.018	-1.447	0.148	-0.063	0.009
month_nov	-0.0620	0.020	-3.036	0.002	-0.102	-0.022
month_oct	0.2339	0.034	6.804	0.000	0.166	0.301
month_sep	0.1071	0.041	2.613	0.009	0.027	0.187
poutcome_other	0.0598	0.023	2.633	0.008	0.015	0.104
poutcome_success	0.4290	0.027	15.640	0.000	0.375	0.483
poutcome_unknown	-0.0020	0.022	-0.091	0.928	-0.046	0.042

```
=====
```

```
In [21]: plt.plot(X.index,model.predict(X),'*')
```

```
Out[21]: [<matplotlib.lines.Line2D at 0x296069142c8>]
```



```
In [22]: features = model.pvalues.sort_values(ascending = True)
features[features < 0.025]
```

```
Out[22]: duration          2.669761e-212
poutcome_success        3.859958e-55
month_oct               1.015051e-11
contact_unknown         1.990262e-08
month_mar               4.873139e-08
month_jan               5.431500e-04
day                    1.601168e-03
month_nov               2.399081e-03
month_jul               2.700422e-03
loan_yes                4.142675e-03
marital_married         6.180763e-03
poutcome_other          8.452288e-03
month_sep               8.965575e-03
job_retired             2.201513e-02
dtype: float64
```

```
In [23]: selected_features = list(features[features<0.025].index) + ['y_yes']
print(selected_features)
```

```
['duration', 'poutcome_success', 'month_oct', 'contact_unknown', 'month_mar', 'month_jan', 'day', 'month_nov', 'month_jul', 'loan_yes', 'marital_married', 'poutcome_other', 'month_sep', 'job_retired', 'y_yes']
```

```
In [24]: features_data = df_dum[selected_features]
features_data.head()
```

```
Out[24]:
```

	duration	poutcome_success	month_oct	contact_unknown	month_mar	month_jan	day	month_nov	month_jul	loan_yes	marital_married	poutcome
0	79	0	1	0	0	0	19	0	0	0	1	1
1	220	0	0	0	0	0	11	0	0	1	1	1
2	185	0	0	0	0	0	16	0	0	0	0	0
3	199	0	0	1	0	0	3	0	0	1	1	1
4	226	0	0	1	0	0	5	0	0	0	0	1

```
In [25]: #save to csv
features_data.to_csv('bank_features.csv',index=False)
```

```
In [ ]:
```

